

Wildfire Containment Squad — Study Guide

Sep 27, 2026 · @Adithya Ajay

How to use this guide

Read sections 2 to 6 once end to end: they explain what the system is and how one simulation step runs. Then study sections 7 to 12 with the code open beside you, and rehearse sections 13 to 15 out loud before the viva.

If you have	Read	Goal
15 minutes	2, 11, 16	Say what the project is and what it found, with numbers
1 hour	2 to 6, 11, 14	Explain every agent, the step order and the results; answer the common questions
An evening	Everything, with the code open	Defend every design decision and walk through any file on request

Three things to memorise word for word:

1. **The one-line pitch:** scouts find the fire, a Coordinator auctions fire zones, firefighters bid and fight with risk-aware A*; on the same fires, the coordinated team saves up to 25 percentage points more forest than independent agents.
2. **The step order:** scouts, then Coordinator, then firefighters in id order, then fire spread, then the safety check, then metrics.
3. **The honest finding:** the shared map and coordinated targeting give most of the gain; the auction beats greedy allocation only when several fires compete (Scenario 3).


Every number in this guide comes from `wildfire_squad/results/` and `results/studies/` (720 main runs + 1,090 study runs, 27 September 2026). If you rerun the experiments, the numbers stay the same, because every run is seeded.

The project in one page

Wildfire Containment Squad is a Python simulation (Mesa 3.5) in which 3 Scout Drones, 5 Firefighters and 1 Coordinator contain a wind-driven wildfire on a 50×50 forest grid. On the same fires, the coordinated team saves 4 to 25 percentage points more forest than independent agents in 7 of 8 test configurations.

The problem. A wildfire spreads faster than one responder can act, smoke limits what anyone can see, and wind changes where the fire runs. Containment needs three things at once: finding the fire, deciding who fights which part, and moving there safely. No single agent can see or reach the whole fire, which is why it is a multi-agent problem.

Scouts sense, the Coordinator allocates, firefighters act on the fire



system architecture · 3 agent types, 1 shared map, 1 environment

Read the loop from the bottom left. The environment gives each agent only what is near it, scouts write that into the shared belief map, the Coordinator turns the map into fire zones and auctions them, firefighters bid, travel and act, and their actions change the environment.

What was built (folder `wildfire_squad/`, about 3,300 lines of Python plus 900 lines of tests):

- A stochastic fire model with the exact Review 1 formula and parameters.
- Three agent types with their own sensing, memory and decision rules, talking only through messages and the shared map.
- Three strategies on identical fires: **independent** (no sharing), **greedy nearest** (shared map, nearest firefighter per zone) and **action** (the full system).
- A live browser demo, 114 automated tests, 720 experiment runs over the six Review 1 scenarios and 1,090 runs of supporting studies.

Headline results (30 seeds per scenario, paired on the same fires):

- The full system beats independent agents significantly in 7 of 8 configurations: +13.5 points with strong wind, +11.1 with the river, +25.3 with 10 firefighters.
- The gain grows with team size: +4.1, +13.5 and +25.3 points for 2, 5 and 10 firefighters.
- The auction beats greedy allocation significantly only with 3 simultaneous fires (+1.2 ± 0.9), exactly the case Review 1 argued for; elsewhere they tie.

- Risk-aware A* expands 63% fewer nodes than uniform-cost search for the same optimal paths.
- Firefighters are caught in about 0.1 cases per run.

The rubric and where each mark is earned

Each of the eight criteria (20 marks in total) has concrete evidence you can point to. Review 1 is the design report; Review 2 is this implementation.

Review	Criterion	Marks	Evidence to point at	What to say
1	PEAS formulation	3	Review 1 §3: a PEAS table for each agent + one shared score formula	Each agent perceives and acts differently, so each has its own PEAS; all three share one score, so their goals pull the same way
1	Environment & agent analysis	3	Review 1 §4–5: 7 environment properties with justification; architecture table with "why not simpler"	Partially observable + stochastic + dynamic + multi-agent is the hardest combination; every design choice answers one of them
1	Algorithmic modeling & search strategy	3	Review 1 §6–7: formal problem formulation, A* admissibility argument, auction formula, comparison tables	A* is the only candidate that is optimal with risk costs and fast enough to replan; the auction is decentralised and cheap to rerun

Review	Criterion	Marks	Evidence to point at	What to say
1	Q&A & presentation	1	Review 1 §10 + section 14 of this guide	Short, confident answers with a number or a test behind each
2	Tool/package selection & setup	3	README §1: tool table with reasons, pinned <code>requirements.txt</code> , <code>setup.bat</code> / <code>setup.sh</code> , clean installs verified on Python 3.12 and 3.13	Mesa is built for agent-based models; NumPy vectorises the fire; setup is one command
2	Multi-agent execution & interaction	3	<code>messages.py</code> (7 message types on one bus), shared belief map, auction log and message counts in the demo, <code>test_auction.py</code> , <code>test_coordination.py</code>	Agents genuinely depend on each other: scouts feed the map, the Coordinator allocates from it, firefighters bid and report back
2	Demo quality & testing scenarios	3	Live demo with belief-map toggle, Compare button and Coordinator-failure switch; 114 tests at 89% coverage; 6 scenarios × 3 strategies × 30 seeds + 4 studies	Every claim has a test or a chart; comparisons are paired on identical fires
2	Code structure & scalability	1	Package layout, Mesa-free <code>algorithms/</code> , everything in <code>config.yaml</code> , <code>test_scales_to_any_agent_count</code> (0-30 agents)	Change a number in the config, not the code

Where marks are most at risk, and the answer ready for each:

- *"Your results don't fully match Review 1's prediction."* Section 11 has the honest comparison: confirmed for Scenario 2 and team scaling, partly for Scenario 3, not for auction vs greedy in general, with the reason.
- *"Why do agents act 5 times per fire update when Review 1 says every step?"* Section 12.1: at the literal timing even perfect knowledge saves only 15%, so no strategy can be compared.
- *"You changed the plan's firebreak rule."* Section 12.2: Review 1 itself calls the firefighter utility-based; the utility rule saves 8 points more, with a 95% confidence interval that excludes zero.

The environment

The world is a 50×50 grid of cells indexed (x, y) , where a wind-driven fire spreads by chance every 5 agent steps. It is partially observable, stochastic, sequential, dynamic, discrete and multi-agent.

Cells and the map

State	Value	Can burn?	Firefighter can enter?
Empty	0	No	Yes
Tree	1	Yes (fuel 1.0 dense, 0.5 sparse)	Yes
Burning	2	Already burning	No
Burnt	3	No	Yes
Firebreak	4	No	Yes
Water	5	No	No (refill next to it)
Unknown	-1	Belief map only	Treated as passable

The random map has 2 lakes, 85% of land is Tree, and 30% of trees sit in rectangular sparse patches (the light-green squares in the demo). All agents start at the base station $(2, 2)$. Ignitions are Tree cells at least 15 cells from the base. The **river** preset puts a water column at $x = 25$ with two 3-cell crossings ($y = 10-12$ and $38-40$) and the fire on the far bank. Code: `environment/forest.py`.

How fire spreads

Each fire update, a Tree cell n next to a Burning cell b ignites with this probability:

$$P(\text{ignite from } b) = p_{\text{base}} \times \text{fuel}(n) \times \left(1 + k \cos \theta\right)$$

Here $p_{\text{base}} = 0.3$, $k = 0.8$ and θ is the angle between the wind and the direction from b to n . Several burning neighbours act as independent chances:

$$P(\text{ignite}) = 1 - \prod_b \left(1 - P(\text{ignite from } b)\right)$$

A burning cell becomes Burnt after 4 fire updates. Empty, Water, Burnt and Firebreak cells never ignite. Code: `environment/fire.py` (`ignition_probability`, `spread_step`, vectorised with NumPy).

Worked example (wind blowing east, dense fuel):

Neighbour direction from the fire	$\cos \theta$	P per update	P within 4 updates
Downwind (east)	1.00	$0.3 \times 1.8 = 0.54$	0.96
Diagonal downwind	0.71	0.47	0.92
Crosswind (north/south)	0.00	0.30	0.76
Diagonal upwind	-0.71	0.13	0.43
Upwind (west)	-1.00	$0.3 \times 0.2 = 0.06$	0.22

Sparse fuel halves every value in the "per update" column. With two burning neighbours, one downwind (0.54) and one crosswind (0.30), the cell ignites with $1 - 0.46 \times 0.70 = 0.68$. `test_fire.py` checks the $1.8\times$ and $0.2\times$ factors exactly.

Time scale

Agents act every step; the fire updates every 5 steps (`fire.spread_every: 5`). Every Review 1 probability is unchanged; agents simply get 5 actions per fire update. At 1 (the literal timing) the fire burns about 90% of the forest in roughly 100 steps and no strategy, not even one with perfect knowledge, saves more than 15% (section 12.1).

When a run ends and how it is scored

A run ends when no cell is burning (contained) or at $t_{\text{max}} = 300$ steps. The Review 1 score is:

$$\text{Score} = 0.7 \cdot \frac{T_{\text{saved}}}{T_{\text{initial}}} - 0.2 \cdot \frac{t_{\text{contain}}}{t_{\text{max}}} - 0.1 \cdot \frac{A_{\text{lost}}}{A_{\text{total}}}$$

T is Tree cells (firebreaks do not count as saved), t is steps (t_max if never contained) and A is firefighters. The primary metric in every chart is **% forest saved** = trees standing at the end ÷ trees at the start × 100. Code: `WildfireModel.score`, `final_metrics`.

The three agents

Scouts perceive, the Coordinator decides who goes where, and firefighters act; none of them can contain the fire alone. All three share one performance measure, the score in section 4.

PEAS

Agent	Performance	Environment	Actuators	Sensors
Scout Drone (3)	Fire found early; grid explored; low belief-map staleness	Forest seen from above; cannot be harmed by fire	Fly 1 cell in 8 directions or hover; broadcast OBSERVE and TARGET	Camera radius 3 (cell states); wind sensor; own position
Firefighter (5)	Cells extinguished; firebreaks cut; zones contained; never caught	Forest on the ground; fire, water, other firefighters; 10 units of water	Move 1 cell in 4 directions; extinguish an adjacent burning cell; cut a firebreak; refill next to water; send BID and DONE	View radius 1; the shared map; own position and water; ANNOUNCE, AWARD, REVOKE
Coordinator (1)	Short travel for assignments; every zone covered; few idle firefighters	The shared map and the firefighter team; no body on the grid	Send ANNOUNCE, AWARD, REVOKE	The shared map; BID and DONE messages

Architecture and memory

Agent	Architecture	Internal state (in code)	Why not simpler
Scout Drone	Model-based, goal-based	The shared belief map; its current target and when it chose it	A reflex scout circles the fire it already sees and never finds new ignitions

Agent	Architecture	Internal state (in code)	Why not simpler
Firefighter	Model-based, utility-based	Position, water, assigned zone, target cell, A* path, steps since planning, when it started refilling	A goal-based firefighter treats every burning cell as equally good; utility prefers cells that stop the most spread per step of travel
Coordinator	Utility-based	Zone table (id, cells, threat, reference cell), assignments firefighter → zone	It trades off distance, threat and water across the whole team, which needs a number

Scout decision cycle (`agents/scout.py`)

1. **Sense** every cell within radius 3 and send one OBSERVE to the belief map.
2. **Choose a goal** only when it has none, has reached it, or someone has seen it since it was chosen: the frontier cell with the highest value = $\text{age} \div (1 + \text{distance})$, skipping cells within 5 of another scout's target.
3. **Publish** the target in a TARGET message so other scouts avoid it.
4. **Fly** one cell straight toward it (it flies, so no path search is needed).

Firefighter decision cycle (`agents/firefighter.py`)

Every step it senses radius 1, reports DONE(contained) if nothing of its zone still burns, then takes **exactly one** action by priority:

1. **Safety first:** standing on a Tree with fire in any of the 8 neighbours → cut its own cell into a firebreak.
2. **Out of water** → walk to the nearest water-side cell; refill to 10 when next to water.
3. **Burning cell next to it** (4 directions) → extinguish the one with the highest downwind spread probability (costs 1 water).
4. **Standing on its firebreak target** and it is a Tree → cut a firebreak.
5. **Otherwise** → take one step of the risk-aware A* path toward its target.

Where the target comes from depends on its situation:

Situation	Target
Holds a zone (auction or greedy)	The zone's burning cell with the highest spread threatened $\div (1 + \text{distance})$, skipping cells a teammate targets; if none, the zone's downwind edge
Free, no fire known anywhere	A frontier cell to explore, like a scout
Free, fire known, waiting for an AWARD	Holds position (the Coordinator usually awards it within a step)
No AWARD for 10 steps (Coordinator silent)	Fallback: the nearest known burning cell
Independent strategy	The nearest burning cell in its private map, else explore

When bidding it computes $U(f, z)$ (section 7) with a real A* path cost, so an unreachable zone gets no bid.


Coordinator decision cycle (`agents/coordinator.py`)

1. **Every 5 steps:** cluster believed-burning cells into zones (8-connected components), keep zone ids stable by maximum overlap, and score each zone's threat.
2. **Release:** take zones back (REVOKE) from firefighters that were lost, refilled for more than 15 steps, or whose zone has disappeared.
3. **Allocate:** fill each zone's slots (1, or 2 for zones above 20 cells) by sequential auction, highest threat first; then give each still-free firefighter one extra slot so nobody idles.
4. **Between rounds:** as soon as a firefighter becomes free, allocate it immediately (event-driven).
5. **Failure test:** from `coordinator.fail_at_step` it does nothing at all.

One simulation step, the messages, and an auction round

`WildfireModel.step()` in `model.py` runs the same six stages every step, in a fixed order, so every run is reproducible and every conflict has a defined winner.

Every step runs the same six stages in a fixed order



one simulation step · `WildfireModel.step()` · 6 stages

Firefighters act before the fire spreads, so a firefighter that cuts its own cell to a firebreak in stage 3 is safe in stage 4. The fire only moves on steps 5, 10, 15 and so on; on the other four steps the world changes only through the agents.

The message protocol (`messages.py`)

Every interaction goes through one `MessageBus` with instant delivery. It counts each type, which is what the demo shows as "Messages".

Message	From → to	Carries	When
OBSERVE	Scout or firefighter → belief map	Cell states seen + the step	Every step, every agent (firefighters only with the shared map)
TARGET	Scout (or exploring firefighter) → the others	The frontier cell it heads for	When it picks a new target
ANNOUNCE	Coordinator → each free firefighter	Zone id, cells, threat, reference cell	For every open slot of a zone, in the auction
BID	Firefighter → Coordinator	Zone id, utility $U(f, z)$	In reply to ANNOUNCE, only if $U > 0$
AWARD	Coordinator → winner	Zone id, reference cell	After the bids for that slot close
DONE	Firefighter → Coordinator	Zone id, <code>contained</code> or <code>abandoned</code>	Zone no longer burns, or no safe path exists

Message	From → to	Carries	When
REVOKE	Coordinator → firefighter	Zone id, reason (absent , zone_gone)	Refilling for more than 15 steps, or the zone disappeared (the one type added to Review 1)

A worked auction round (illustrative numbers)

The Coordinator knows two zones: **Z1** has 30 burning cells and threat 4.2; **Z2** has 8 cells and threat 1.5. Three firefighters are free. Each bids $U = \text{threat} \div (1 + A * \text{path cost}) \times \text{water} \div 10$.

Firefighter	Water	Path cost to Z1	Path cost to Z2	Bid for Z1	Bid for Z2
FF5	10	14	6	$4.2 \div 15 \times 1.0 = \mathbf{0.280}$	$1.5 \div 7 \times 1.0 = 0.214$
FF6	4	9	20	$4.2 \div 10 \times 0.4 = 0.168$	$1.5 \div 21 \times 0.4 = 0.029$
FF7	10	25	10	$4.2 \div 26 \times 1.0 = 0.162$	$1.5 \div 11 \times 1.0 = 0.136$

1. Z1 goes first because its threat is higher, and it gets **2 slots** because it has more than 20 cells.
2. Slot 1: ANNOUNCE to FF5, FF6, FF7 → bids 0.280, 0.168, 0.162 → **AWARD to FF5**.
3. Slot 2: FF5 has left the pool → bids 0.168, 0.162 → **AWARD to FF6**.
4. Z2, 1 slot: only FF7 is left → bid 0.136 → **AWARD to FF7**.

Greedy nearest would take zones in discovery order, ignore threat and water, and could send full-water FF5 to the small fire while half-empty FF6 faces the large one. Ties go to the lower id (`pick_winner`); a firefighter that cannot reach a zone bids nothing. After these design slots, any firefighter still free gets one extra "support" slot so nobody idles.

Algorithms in depth

Five algorithms do the work: a timestamped belief map (perception), frontier exploration (scouts), connected-component zones (Coordinator), risk-aware A* (navigation) and a sequential single-item auction (allocation). All live in `algorithms/` or `belief_map.py` as plain functions with no Mesa dependency, so each is unit-tested on its own.

1. Shared belief map (`belief_map.py`)

For every cell it stores the last observed state ($-1 = \text{unknown}$) and `seen_step`, the step it was observed.

- **Fusion rule:** an observation is written only if its step $\geq$ the stored step, so the **newest timestamp wins** and a late, older report is ignored.
- **Age** = current step - `seen_step`; a cell is **stale** when its age exceeds 15.
- **Terrain prior:** agents know where water is (a map of the forest), not where the fire is.
- In the independent strategy each firefighter has its **own private** map and never sees scout data.

2. Frontier exploration (`algorithms/frontier.py`)

Candidates are unknown or stale cells; the **frontier** is the candidates next to a known, fresh cell. A scout at position p picks:

$$\text{target} = \arg\max_{c \in \text{frontier}} \frac{\text{age}(c)}{1 + \text{chebyshev}(p, c)}, \quad \text{age}(\text{unknown}) = t_{\max} + 1 = 301$$

Cells within 5 (Chebyshev) of another scout's published target are skipped. If that rules out the whole frontier (early on, when it is a small ring round the base), every unknown or stale cell becomes a candidate. The scout re-picks only when it reaches the target or anyone observes it. Three scouts explore more than 80% of the map (`test_scouts.py`).

3. Fire zones and threat (`algorithms/zones.py`)

- **Clustering:** believed-burning cells are grouped into 8-connected components by breadth-first search.
- **Stable ids:** each old zone's id goes to the new component that overlaps it most. A merge keeps the id with the largest overlap; in a split the largest piece keeps the id and the rest get new ids; a zone with no successor is removed (its firefighters get REVOKE).
- **Live zones:** between clustering rounds a firefighter tracks its zone as the burning cells connected to the last snapshot, so it never chases a cell that has already burnt out.

$$\text{threat}(z) = |z| \times \operatorname{mean}_{n \in N(z)} \max_{b \in z, b \sim n} P(\text{ignite } n \text{ from } b)$$

$N(z)$ is the believed-Tree cells touching the zone. A big zone with wind pushing it into dense forest scores highest. Each zone also gets a **reference cell** on its downwind edge, the cell used to price bids.

4. Risk-aware A* (`algorithms/astar.py`)

The search runs on a 4-connected grid; Burning and Water cells cannot be entered. Entering cell n costs:

$$c(n) = 1 + \lambda \cdot \text{risk}(n), \quad \text{risk}(n) = \max_{b \text{ burning next to } n} P(\text{ignite } n \text{ from } b), \quad \lambda = 5$$

So a step costs 1 far from fire and up to $1 + 5 \times 0.54 = 3.7$ right downwind of it. The heuristic is Manhattan distance to the nearest goal cell (the target, or its free neighbours if the target itself is burning):

$$h(n) = \min_{g \in \text{goals}} (|x_n - x_g| + |y_n - y_g|)$$

- **Admissible:** every move changes x or y by exactly 1 and costs at least 1, so the true remaining cost is at least the number of moves, which is at least $h(n)$. A* therefore returns an optimal path.
- **Consistent:** h changes by at most 1 per move while each move costs at least 1, so $h(n) \leq c(n \rightarrow n') + h(n')$ and no node is expanded twice.
- **Implementation:** `heapq` ordered by $(f, h, \text{counter}) \rightarrow$ ties prefer nodes closer to the goal, then insertion order; a closed set; it returns `(path, cost, nodes_expanded)`.
- **UCS** is the same function with $h = 0$. Every A* query in the experiments is repeated with UCS for measurement only: same cost, 63% fewer nodes for A*.
- **Replanning:** when the next cell is burning or blocked, every 10 steps, or when the target changes; no path $\rightarrow$ DONE(abandoned) and the zone is re-auctioned.
- **Reservations** block only the first move (other firefighters' cells this step). If the only way on is through a reserved cell, the firefighter keeps that path and waits.

5. Sequential single-item auction (`algorithms/auction.py`)

$$U(f, z) = \frac{\text{threat}(z)}{1 + \text{pathcost}(f, z)} \times \frac{\text{water}_f}{\text{water}_{\max}}$$

- Zones are auctioned in decreasing threat; a zone above 20 cells has 2 slots.
- For each slot: ANNOUNCE to every free firefighter, collect positive BIDs, AWARD the highest (tie $\rightarrow$ lower id), remove the winner from the pool, repeat.
- Unreachable (infinite path cost) or no water $\rightarrow U = 0 \rightarrow$ no bid.
- **Cost:** about $Z \times F$ A* calls per round (Z zones, F free firefighters, both usually under 10). The Hungarian algorithm would be optimal for one snapshot but is $O(n^3)$, fully central, and must be recomputed every time zones change.
- **Greedy baseline** (`greedy_assignment`): zones in discovery order, each slot to the nearest free firefighter by straight-line distance; ignores threat and water.

6. Choosing a cell inside the zone

Once assigned, a firefighter picks the zone's burning cell b with the highest:

$$\text{value}(b) = \frac{\sum_{n \text{ fuel next to } b} P(\text{ignite } n \text{ from } b)}{1 + \text{manhattan}(\text{firefighter}, b)}$$

This means "spread stopped per step of travel". Cells a teammate already targets are skipped. It is the utility-based behaviour Review 1 describes, and it saves 8 points more than the plan's downwind-edge rule (section 12.2).

Conflicts and how each is resolved

Every conflict that can arise between agents has one deterministic rule and a test that checks it. The first six are the Review 1 conflict table; the last three came up during implementation.

Conflict	Example	Rule	Code	Test
Two firefighters want the same zone	Both are close to the biggest fire	Auction: higher utility wins, tie → lower id; the loser bids on the next zone	<code>sequential_auction, pick_winner</code>	<code>test_zones_auctioned_by_threat_and_loser_bids_on_next</code>
Two firefighters want the same cell in one step	Paths cross in a narrow gap	They act in id order; each current and newly entered cell is reserved; a later firefighter waits, and replans after 2 blocked steps	<code>WildfireModel.step, Firefighter._go_to</code>	<code>test_two_firefighters_never_share_a_cell, corridor test</code>
Two scouts head for the same area	Both rate one stale region highest	Published TARGET messages; cells within 5 of another scout's target are excluded	<code>select_target</code>	<code>test_scout_targets_respect_exclusion_radii</code>
Conflicting observations of one cell	Scout saw it burning at step 40; a firefighter saw it burnt at step 42	Newest timestamp wins	<code>BeliefMap.observe</code>	<code>test_newest_observation_wins, test_older_observation_is_ignored</code>
Path blocked by new fire	The front crosses a planned route	Replan with A*; no safe path → DONE(abandoned) and the zone is re-announced	<code>_go_to, _send_done</code>	<code>test_firefighter_replans_when_path_blocked_by_fire, test_abandoned_zone_is_offered_again</code>
Out of water mid-task	Water hits 0 inside the zone	Go to refill and keep the zone; after 15 steps away the Coordinator REVOKES and re-auctions it	<code>_refill, release_absent</code>	<code>test_absent_firefighter_is_reauctioned</code>
An idle firefighter blocks the only route	The one at the base is walled in by 4 teammates	The blocked one asks; the idle one steps aside to a free cell	<code>request_yield, yield_if_asked</code>	<code>test_boxed_in_firefighter_is_not_deadlocked</code>

Conflict	Example	Rule	Code	Test
Zones merge or split	Two fronts join	Merge keeps the id with the largest overlap; a split's largest piece keeps the id; firefighters of a vanished zone get REVOKE	<code>match_zones</code> , <code>recluster</code>	<code>test_merge...</code> , <code>test_split...</code>
The Coordinator goes silent	It fails at step 60	After 10 steps without an AWARD, each firefighter attacks the nearest known fire on its own	<code>_await_award_or_fal</code> <code>lback</code>	<code>test_coordinator_fail</code> <code>ure_triggers_fallback</code>

The first rule is **negotiation**, the second and third are **reservations of shared resources**, the fourth is **belief fusion**, and the rest are **replanning and recovery**. That covers the kinds of interaction the "Multi-Agent Execution & Interaction" criterion lists: communication, negotiation, cooperation and conflict resolution.

Code tour

The code splits into four layers: the environment, pure algorithms, agents, and the model that runs them; the demo and experiment scripts sit on top. Read it in the order of this table.

Order	File	Lines	What to find there
1	<code>config.yaml</code>	85	Every parameter; values marked <code>impl</code> are implementation choices, the rest are Review 1 values
2	<code>environment/cells.py</code> , <code>fire.py</code> , <code>forest.py</code>	354	Cell states; <code>ignition_probability</code> , <code>spread_step</code> , <code>risk_map</code> ; map generation and presets
3	<code>belief_map.py</code>	71	<code>observe</code> (newest wins), <code>age</code> , <code>stale_mask</code> , <code>passable_mask</code>
4	<code>algorithms/astar.py</code>	142	<code>astar</code> , <code>ucs</code> , <code>plan_path</code> , <code>goal_cells</code>
5	<code>algorithms/frontier.py</code> , <code>zones.py</code> , <code>auction.py</code>	442	<code>select_target</code> ; <code>cluster_zones</code> , <code>match_zones</code> , <code>zone_threat</code> , <code>live_zone_cells</code> ; <code>utility</code> , <code>sequential_auction</code> , <code>greedy_assignment</code>

Order	File	Lines	What to find there
6	<code>messages.py</code>	120	The 7 message dataclasses and <code>MessageBus</code>
7	<code>agents/scout.py</code>	87	<code>sense</code> , <code>choose_target</code> , <code>move</code>
8	<code>agents/coordinator.py</code>	175	<code>recluster</code> , <code>release_absent</code> , <code>allocate</code> , <code>_collect_bids</code>
9	<code>agents/firefighter.py</code>	490	<code>step</code> (action priority), <code>receive</code> (messages), <code>bid_utility</code> , <code>_choose_zone_target</code> , <code>_go_to</code> (navigation + reservations), <code>fallback</code>
10	<code>model.py</code>	268	<code>WildfireModel.__init__</code> , <code>step</code> , <code>plan</code> (A* + UCS measurement), metrics and <code>score</code>
11	<code>strategies.py</code> , <code>scenarios.py</code> , <code>settings.py</code>	123	The three strategies' switches; the six scenarios; config loading and deep-merge
12	<code>app.py</code> , <code>run_experiments.py</code> , <code>run_studies.py</code>	1,024	Demo; main experiment; supporting studies

How a run is wired together

1. `WildfireModel(config_overrides, seed)` loads `config.yaml`, deep-merges the overrides (a scenario, a strategy) and seeds Mesa.
2. It creates **three independent random streams** from the seed: one for the map, one for the fire, and Mesa's own for agents. The map and fire are therefore identical for every strategy on the same seed.
3. `build_forest` makes the map and ignitions; the model sets the ignitions burning and creates the `MultiGrid`.
4. It creates the `MessageBus` and the shared `BeliefMap`, and subscribes handlers: `OBSERVE` → the map, `TARGET` → the scout and explorer target tables.
5. It creates scouts at the base, a Coordinator (not for independent), and firefighters on distinct base cells, each with the shared map or a private one.
6. `step()` runs the six stages; `run()` loops until contained or step 300 and returns `final_metrics()`.

Design patterns to name if asked

- **Blackboard:** the shared belief map is a common workspace that agents read and write.
- **Contract net:** ANNOUNCE → BID → AWARD is the classic contract-net protocol for task allocation.
- **Strategy switch:** one code path; `strategies.py` turns sharing and the Coordinator on or off.
- **Pure core, thin shell:** algorithms take arrays and callbacks, so tests call them directly without a simulation.
- **Caching:** `WildfireModel.plan` caches the passable mask and cost grid per belief-map version, so many A* calls in one step reuse them.

Strategies, scenarios and how testing is organised

Testing has three layers: 114 automated tests prove each rule works; 720 experiment runs compare the three strategies on the six Review 1 scenarios; 1,090 study runs test the design decisions.

The three strategies (`strategies.py`)

Strategy	Shared map?	Coordinator?	Allocation	Role in the study
Independent	No: each firefighter has a private map (scouts still fly, but nobody reads their map)	No	Each goes to the nearest fire it has seen, or explores	Baseline 1: "no multi-agent coordination"
Greedy nearest	Yes	Yes	Zones in discovery order; each slot to the nearest free firefighter	Baseline 2: coordination without the auction
Auction	Yes	Yes	Sequential auction on $U(f, z)$, highest threat first	The full system

All three share the same firefighter code, A* and safety rules, so a difference in results comes from sharing and allocation alone.

The six Review 1 scenarios (`scenarios.py`)

#	Name	Change from default	What it tests
1	<code>calm_single</code>	Wind $k = 0$, 1 ignition	Baseline behaviour; fire grows as a ring
2	<code>windy_single</code>	Wind $k = 0.8$, 1 ignition (= default)	Fast directional spread
3	<code>multi_ignition</code>	$k = 0.4$, 3 ignitions ≥ 15 cells apart	Allocation across several zones
4	<code>no_scouts</code>	0 scouts	Value of shared perception
5	<code>scaling_2 / _5 / _10</code>	2, 5 or 10 firefighters	Scalability and diminishing returns
6	<code>river</code>	Water column at $x = 25$, two 3-cell crossings, fire on the far bank	Path conflicts and replanning

One extra map, **competing**, is used only in the allocation study: a slow fire in sparse forest near the base plus a fast fire in dense forest farther away.

How the experiment is kept fair

- **Paired seeds:** seeds 0–29, and the same seed gives every strategy the same map and the same fire random numbers.
- **Paired statistics:** for each seed, auction minus baseline; the mean and a 95% confidence interval ($1.96 \times \text{std} \div \sqrt{30}$) are reported, plus how many of the 30 seeds the auction won.
- **Deterministic:** the same seed and config always give the same result (`test_same_seed_same_result`).

The automated tests (`tests/` , 114 tests, about 40 s)

File	Tests	Proves
<code>test_astar.py</code>	31	A* cost = UCS cost on 125 random queries, never more nodes, avoids fire, detours at high λ , no path when blocked
<code>test_model_smoke.py</code>	18	Every strategy and scenario runs, 0–30 agents, score formula, determinism


File	Tests	Proves
<code>test_coordination.py</code>	12	Re-auction, abandon, event-driven allocation, Coordinator failure + fallback, safety, time scale, presets
<code>test_fire.py</code>	9	The Review 1 fire formula, clipping, burn-out, determinism, ignition placement, river layout
<code>test_auction.py</code>	9	Utility, highest bid wins, tie-break, threat order, 2 slots for big zones, no bid when unreachable, full message flow
<code>test_zones.py</code>	8	Clustering, stable ids, merge/split, threat, downwind targets
<code>test_tools.py</code>	8	Demo helpers, comparison, experiment and study runners end to end
<code>test_conflicts.py</code>	7	No shared cells, corridor waiting, scout exclusion, replanning, no deadlock
<code>test_scouts.py</code>	7	Frontier, exploration > 80%, shared vs private maps
<code>test_belief_map.py</code>	5	Newest wins, older ignored, staleness, passability

Coverage is 89% of all code; the simulation modules are 93–100% each. The uncovered part is mostly the demo's page layout, which was checked in a real browser instead.


Results and what they mean

The full system saves the most forest in 6 of 8 configurations and beats independent agents significantly in 7 of 8. Against greedy allocation it wins significantly only with three simultaneous fires.

Coordinated teams save the most forest



Auction gain on the same fires



Dot = mean paired gain over 30 seeds; whisker = 95% confidence interval. Above 0 = auction better.

results/summary.csv and results/paired_gains.csv · 8 configurations × 3 strategies × 30 seeds

Read the second chart like this: a whisker that stays above the grey zero line means the auction is better on the same fires with 95% confidence. Blue (vs independent) clears zero everywhere except three fires; orange (vs greedy) clears it only for three fires.

Why the bars in the first chart vary so much between scenarios: some ignitions start in sparse forest and burn out early; others run across the map. That is why every comparison is paired on the same seed; the paired gains remove that variance. S2 and S5-with-5-firefighters are the same default configuration, so their numbers are identical, which also shows the runs are deterministic.

What each result means

- **Shared perception matters most.** Independent firefighters see 1 cell around them, so they find the fire late and crowd it. With 10 of them they save only 44%; the coordinated team saves 69%.
- **The gain grows with team size** (+4, +13, +25 points for 2, 5, 10 firefighters): extra firefighters are only useful if someone tells them where to go.
- **Without scouts the team still wins** (+6.9): firefighters share one map and publish their exploration targets, so five of them search like one coordinated sweep.
- **The river** is where A* shows its value: independent firefighters never extinguish a single cell on the far bank; coordinated ones path through the crossings and win on 29 of 30 seeds.
- **Three fires overwhelm five firefighters** under every strategy (at most 6% saved). The auction still beats greedy there, because it sends crews to the most threatening zone first.

Findings vs the Review 1 hypothesis

Review 1 predicted: *"The full system should save the most forest, with the largest margin in Scenarios 2 and 3."*

Prediction	Verdict	Evidence
The full system saves the most	Confirmed	Highest mean in 6 of 8; significantly above independent in 7 of 8
Largest margin in Scenario 2	Confirmed	+13.5 over independent, among the largest single-fire margins
Largest margin in Scenario 3	Partly	The only scenario where the auction beats greedy significantly (+1.2), but three fires overwhelm five firefighters, so the margin over independent is small
The auction beats greedy in general	Not confirmed	Tied within noise elsewhere: with one or two zones there is rarely a real choice to make

Say it in one sentence: *"Most of the benefit comes from the shared belief map and coordinated targeting; the auction's threat-aware allocation adds a measurable benefit only when several fires compete."*


Other measured facts

- **A* vs UCS:** 63% fewer nodes expanded overall (49–87% per scenario) for the same optimal cost.
- **Safety:** 0.10 firefighters caught per run with the auction (0.16 greedy, 0.13 independent).
- **Work done:** about 75 cells extinguished per run by coordinated crews vs 29 by independent ones.
- **Speed:** about 1.4 s per 300-step run, including the UCS measurement.

Supporting studies and the decisions they justify

Four studies (`run_studies.py` , 1,090 runs) test the choices the implementation had to make. Each decision below has a study or a test behind it, and each is a switch in `config.yaml` , so an evaluator can turn it off and see the difference.

12.1 Fire time scale: agents act 5 times per fire update




Windy scenario, 20 seeds per point. At 1 agent step per fire update even perfect knowledge saves 15%.

results/studies/timing_summary.csv · 5 time scales × 4 variants × 20 seeds

At the literal Review 1 timing (1 on the x axis) all three strategies save the same 11.6% (auction minus independent = $+0.0 \pm 0.7$), and even firefighters told the true fire state save only 15%. The fire crosses the forest in about 100 steps, faster than any team can act, so the literal timing cannot test Review 1's hypothesis at all. As agents get more actions per update, coordination starts to pay: at 5 the auction gains $+14.1 \pm 5.7$ over independent agents. Every probability stays exactly as in Review 1; only the unit of time changes.

12.2 What a firefighter targets inside its zone



The implementation plan sent firefighters to the zone's downwind edge, the head of the fire. That is the most dangerous place and the least effective one, because the head outruns anyone working there. Review 1's own agent analysis says the firefighter is utility-based so it can "prefer cells that stop the most future spread". The utility rule does exactly that and saves $+8.1 \pm 4.2$ points more than the downwind-edge rule (windy, auction), so it is the default. The plan's rule is still available as `zone_tactic: firebreak`.

12.3 Targeted allocation test

A slow fire in sparse forest near the base plus a fast fire in dense forest farther away, with only 3 firefighters: the case where threat-aware allocation should matter most.

Variant	Forest saved
Independent	22.6%
Greedy (8-connected zones)	30.3%
Auction (8-connected zones)	29.5%
Greedy (fragments within 3 cells joined)	33.4%
Auction (fragments within 3 cells joined)	33.3%

Auction and greedy tie (-0.8 ± 2.0); both beat independent ($+6.9 \pm 3.4$). The most likely reason: the near fire is found first, so both methods commit crews to it before the far fire is known. Joining fragments helps both by about 3 points, but that departs from the plan's 8-connectivity, so the default stays at 1.

12.4 Coordinator failure and the fallback

Variant	Forest saved
Independent	42.6%
Auction, Coordinator works	56.1%
Auction, Coordinator fails at step 60, fallback on	60.7%
Auction, Coordinator fails at step 60, no fallback	43.9%

Without the fallback a failed Coordinator drops the team to the independent level. With it (no AWARD for 10 steps → attack the nearest known fire), the team keeps working: $+16.8 \pm 4.9$ over no fallback. It even edges out the working Coordinator here ($+4.6 \pm 4.1$): with one fire, once every crew is engaged, "everyone to the nearest fire on the shared map" is as good as zone allocation. This agrees with section 11: the shared map carries most of the gain.

Every implementation decision at a glance

Decision	Config key	Why	Evidence
Agents act 5× per fire update	<code>fire.spread_every: 5</code>	At 1, nothing can be compared	12.1
Utility targeting inside a zone	<code>firefighter.zone_tactic: utility</code>	Review 1 calls the firefighter utility-based; +8 points	12.2
Live zones + event-driven allocation	built in	Crews chased burnt-out cells and idled up to 5 steps	Review 1 §4 "event-driven replanning"; <code>test_coordination.py</code>
Safety first (cut own cell; no stepping onto fuel near fire on a fire step)	<code>firefighter.secure_footing</code>	Losses fell from about 0.9 to 0.1 per run	<code>test_firefighter_secures_footing_before_anything_else</code>
Fallback after 10 steps without AWARD	<code>firefighter.award_timeout: 10</code>	Review 1's single-point-of-failure answer	12.4

Decision	Config key	Why	Evidence
Support rounds for idle firefighters	<code>coordinator.assign_idle</code>	Coordinator PEAS: "few idle firefighters"; greedy gets the same	Both strategies
Terrain known in advance	built in	Firefighters must know where water is to refill	<code>BeliefMap.water_prior</code>
Idle firefighters explore	<code>firefighter.explore_when_idle</code>	Makes Scenario 4 (no scouts) work	S4 result +6.9
REVOKE message	<code>messages.py</code>	A clean way to take back a zone	<code>test_absent_firefighter_is_reauctioned</code>
Distinct start cells + yielding	built in	Two firefighters never share a cell; no deadlock at base	<code>test_conflicts.py</code>

Running everything and the demo script

On Windows, double-click `setup.bat` once, then `run_demo.bat`; everything else is one command from the `wildfire_squad` folder. Python 3.12 or newer is required (Mesa 3.5 needs it).

Task	Windows	Any system	Time
Install	<code>setup.bat</code>	<code>./setup.sh</code> , or <code>pip install -r requirements.txt</code> in a venv	1-2 min
Live demo	<code>run_demo.bat</code>	<code>solara run app.py</code> → <code>http://localhost:8765</code>	starts in ~5 s
All tests	<code>run_tests.bat</code>	<code>python -m pytest</code>	~40 s
Quick experiment check	<code>run_experiments.bat --quick</code>	<code>python run_experiments.py --quick</code>	~30 s
Full experiment + studies	<code>run_experiments.bat</code>	<code>python run_experiments.py --scenarios all --runs 30</code> then <code>python run_studies.py</code>	~10 min
Rebuild charts only		<code>python run_experiments.py --replot</code> , <code>python run_studies.py --replot</code>	seconds

Before the viva: open the demo on the presenting laptop at least once (the first start is slower), set the zoom so the map and the right panel both fit, and keep `results/paired_gains.png` and `results/studies/timing.png` open in another window as backups.

The demo, screen by screen

- **Left:** Reset, Play/Pause and Step, speed, then the parameters. Parameters apply only after **Reset**.
- **Centre:** the forest, with scouts (cyan circles), firefighters (yellow squares with their id), zone reference cells (white x) and the wind arrow. The checkbox switches to the **belief map**: black = never seen, faded = stale.
- **Right:** a line with the current settings (the sliders do not show numbers), % forest saved, zones with threat and crews, each firefighter's water and activity, the **last auction rounds with every bid and the winner**, message counts, recent events.
- **Bottom:** % saved and burning cells over time, and the **Compare** button.

Five-minute demo script (rehearse it)

1. **Set up:** preset `windy_single`, seed 0, strategy `auction`, then Reset. Say: "Three scouts, five firefighters, one Coordinator; wind blows east."
2. **Play.** Scouts fan out (point at the cyan circles). When the fire is found, the right panel fills with ANNOUNCE/BID/AWARD rounds. Say: "Each firefighter bids threat ÷ (1 + path cost) × water; the highest bid wins."
3. **Pause, tick Show belief map.** Say: "This is what the team knows: black is unexplored, faded is out of date. Firefighters plan on this, not on the truth."
4. **Let it finish:** it is contained at step 195 with 84.7% saved.
5. **Press Compare.** It shows independent 49.5%, greedy 59.5%, auction 84.7% on this exact fire, with the three final maps side by side. Say: "Same map, same fire, same random numbers; only the coordination differs."
6. **Preset `river`, Reset, Play** (optional): independent firefighters never reach the far bank; coordinated ones route through the two crossings with A*.
7. **Coordinator fails at step = 60, Reset, Play** (optional): `fallback` events appear on the right and the crews keep fighting. Say: "This answers the single-point-of-failure question from Review 1."

If anything misbehaves, press Reset; if the page hangs, restart `run_demo.bat`. The charts in `results/` show the same story without the live demo.

Viva question bank

Each answer is short enough to say in about 20 seconds and ends with something you can show. Practise them out loud; when unsure, answer with the number and where it comes from.

The idea

Question	Model answer
Why is this multi-agent and not one smart agent?	No single agent can see or reach the whole fire. Perception (scouts) and action (firefighters) are split, and they depend on each other through the shared map and the auction. The Compare button shows it: +35 points on seed 0.
How do the agents actually interact?	Through 7 message types on one bus (OBSERVE, TARGET, ANNOUNCE, BID, AWARD, DONE, REVOKE) and a shared belief map. The demo counts every message.
Is the Coordinator "cheating", making it centralised?	It only allocates; each firefighter computes its own bid from its own path cost and water, and decides its own actions. If it fails, firefighters fall back to acting alone (section 12.4).
What does each agent optimise?	One shared score: $0.7 \times \text{trees saved} - 0.2 \times \text{time to contain} - 0.1 \times \text{firefighters lost}$, so their goals never conflict.

PEAS and environment

Question	Model answer
Why a separate PEAS per agent?	They sense and act differently: scouts see radius 3 and fly, firefighters see radius 1 and walk, the Coordinator has no body.
Why partially observable?	No agent sees the whole grid; the belief map stores what was seen and when. The belief-map toggle shows the unknown and stale areas.

Question	Model answer
Why stochastic?	Ignition is random, $0.3 \times \text{fuel} \times (1 + 0.8 \cos \theta)$, so the same state can lead to different next states.
Why dynamic and sequential?	The fire moves whether or not agents act, and a firebreak cut now changes what burns later.
Why model-based and not reflex?	A reflex agent has no memory of cells it cannot see, and the fire is almost always partly out of view.
Why do agents act 5 times per fire update?	At the literal timing even perfect knowledge saves 15% and all strategies tie; every Review 1 probability is unchanged (timing chart, section 12.1).

Algorithms

Question	Model answer
Why A* and not BFS?	Step costs differ with fire risk, so BFS (equal costs) is not optimal. DFS is not optimal either, and greedy best-first walks into fire.
Is your heuristic admissible?	Yes. Every move costs at least 1 and changes the Manhattan distance by at most 1, so h never overestimates. It is also consistent. <code>test_astar.py</code> checks A* cost = UCS cost on 125 queries.
What does A* gain over UCS?	63% fewer nodes expanded for identical paths and costs (<code>astar_vs_ucs.png</code>).
What is risk(n)?	The highest chance that a burning neighbour ignites n on the next update; the step cost is $1 + 5 \times \text{risk}$, so paths bend away from the fire front.
What if the fire blocks the path mid-route?	The firefighter replans; with no safe path it sends DONE(abandoned) and the zone is re-auctioned.

Question	Model answer
Why an auction and not the Hungarian algorithm?	Zones change every few steps. The auction is cheap to rerun (about $Z \times F$ path searches), uses each agent's own cost, and needs no central recomputation; Hungarian is $O(n^3)$ and fully central.
Is the auction optimal?	No, sequential single-item auctions are greedy per slot. In exchange they are fast and robust, and in our tests they match or beat greedy nearest.
How does threat work?	Zone size $\times$ the average chance its neighbouring trees ignite. A big zone pushed by wind into dense forest scores highest and is auctioned first.
How do scouts avoid duplicating work?	Each publishes its target; others ignore cells within 5 of it. Value = age $\div$ (1 + distance), so they prefer unknown, near cells.
Two observations disagree. Which wins?	The newest timestamp (<code>BeliefMap.observe</code>).

Implementation, testing, results

Question	Model answer
Why Mesa?	It is built for agent-based models: agents, grid, seeded randomness and data collection, plus Solara for the browser demo. We use the Mesa 3 API throughout.
How do you make the comparison fair?	Paired seeds: the same map and the same fire random numbers for every strategy; we report the per-seed difference with a 95% confidence interval.
Does coordination help?	Yes, significantly in 7 of 8 configurations, by up to 25 points, and the gain grows with team size.
Why doesn't the auction beat greedy everywhere?	With one or two zones there is rarely a real choice; both assign everyone. It wins significantly where fires compete (Scenario 3). We tested this directly (section 12.3).

Question	Model answer
Did your results match Review 1?	The full system saves the most, confirmed; the largest margins are in Scenario 2 and with more firefighters; Scenario 3 only partly, because three fires overwhelm five firefighters.
How do you know it scales?	The same code runs 0 scouts and 1 to 30 firefighters; a test checks it, and Scenario 5 measures 2, 5 and 10.
How is it tested?	114 pytest tests (89% coverage): the fire formula, belief fusion, A* optimality, auction rules, every conflict rule, robustness and the tools.
What did you change from the plan, and why?	Time scale, utility targeting, live zones, safety-first rules, the fallback. Each has a study or test behind it (section 12 table).
Is the Coordinator a single point of failure?	It was; now firefighters fall back after 10 silent steps. Demo switch: "Coordinator fails at step".
Limitations?	2D grid; simplified spread; instant, perfect communication; no terrain slope; agents act 5× per fire update; auction ≈ greedy with few zones.
Future work?	Communication delays and losses, a decentralised auction with no Coordinator at all, terrain and real GIS maps, learning the targeting utility from data.

Where the code differs from the Review 1 report

An examiner comparing the report with the code may find these five differences. Each is deliberate and has an answer.

Review 1 says	The code does	Say
A burning cell becomes Burnt after 4 steps	After 4 fire updates ; the fire updates every 5 agent steps, so 20 agent steps	"The probabilities are unchanged; agents act 5 times per fire update because at the literal timing even perfect knowledge saves only 15% (timing study)."
Cut a firebreak on an adjacent tree cell	Converts the tree it is standing on , after walking onto it	"Same actuator, same one cell per action, no water. Standing on the cell makes it safe the instant it is cut, which is also how a firefighter protects itself next to fire (losses fell from 0.9 to 0.1 per run)."
Goal test: position adjacent to the zone's target cell	Walks to the target cell; if that cell is burning or water, to a free cell next to it	"A burning target is reached from next to it, exactly as the report says; a tree target is stood on so it can be cut. <code>goal_cells</code> in <code>astar.py</code> implements both cases."
Mesa provides the grid and a scheduler	Mesa 3 removed schedulers, so <code>WildfireModel.step()</code> calls the agents in a fixed order itself	"The plan required the Mesa 3 API; a fixed order is also what makes the conflict rules deterministic."
At the zone's downwind edge, cut a firebreak	Goes to the burning cell that stops the most spread per step of travel; the downwind edge is the fallback when nothing is worth putting out	"Review 1 calls the firefighter utility-based so it can prefer cells that stop the most spread; that rule saves 8 points more (tactics study)."

Presentation plan

Plan 14 minutes: 10 minutes of slides around a 4-minute live demo, then questions. Each slide earns a named rubric criterion, and one team member owns each block.

Team: Adithya Ajay (CB.SC.U4CSE23102), Amrith B (CB.SC.U4CSE23105), Gowreesh B (CB.SC.U4CSE23119), Vinaayak Kanagaraj (CB.SC.U4CSE23152). The split below is a suggestion; swap blocks freely, but keep one owner per block.

#	Slide	Time	Speaker	The one sentence to land	Rubric
1	Title	0:15	Adithya	"We built and tested the multi-agent wildfire squad designed in Review 1."	—
2	Recap: problem and architecture	1:00	Adithya	"Scouts sense, the Coordinator allocates, firefighters act; no one can do it alone."	Consistency with Review 1
3	Tools and setup	1:00	Adithya	"Mesa for agents, NumPy for the fire, one-click setup verified on Python 3.12 and 3.13."	Tools & setup (3)
4	How the agents interact	1:30	Amrith	"Six fixed stages per step, seven message types, and a rule for every conflict."	Multi-agent interaction (3)
5	Algorithms in the code	1:30	Amrith	"Risk-aware A* is optimal and expands 63% fewer nodes than UCS; the auction prices each zone by threat, distance and water."	Multi-agent interaction (3)
6	Live demo	4:00	Gowreesh	Follow the demo script in section 13: auction run, belief map, Compare	Demo quality (3)
7	How we tested	1:00	Vinaayak	"114 tests, six scenarios × three strategies × 30 paired seeds."	Testing scenarios (3)
8	Results	1:30	Vinaayak	"Coordination wins significantly in 7 of 8 configurations, by up to 25 points."	Testing scenarios (3)
9	Findings vs Review 1	1:00	Vinaayak	"Confirmed overall; the auction beats greedy only where fires compete."	Honesty, Q&A
10	Decisions backed by studies	1:00	Gowreesh	"Every change from the plan has a study behind it."	Execution (3), Q&A
11	Code structure and scalability	0:30	Amrith	"Pure algorithms, one config file, 0 to 30 agents with no code change."	Code structure (1)
12	Limitations and thanks	0:30	Adithya	"2D, instant communication, simplified fire: all future work."	—

Delivery rules

- **Open with the pitch** (section 1), not with the tools.
- **Numbers, not adjectives:** "+13.5 points on the same fires", not "much better".
- **Name the honest finding before they ask:** auction $\approx$ greedy except with competing fires. Owning it reads as rigour; being caught by it reads as a gap.
- **Hand over by name:** "Amrith will now show how the agents talk to each other."
- **In Q&A, whoever owns the slide answers first;** others add only a number or a pointer ("that is in `test_astar.py`").
- **Backup plan:** if the demo fails, go to the Compare screenshot and the charts in `results/`, then keep talking.

Quick reference

The numbers most likely to be asked for, in one place.

Parameters

Parameter	Value	Parameter	Value
Grid	50 × 50	Scouts / firefighters	3 / 5
Tree density	85% of land	Sparse trees (fuel 0.5)	30% of trees
p_base	0.3	Wind strength k	0.8
Downwind / upwind factor	1.8× / 0.2×	Burn time	4 fire updates
Fire updates	every 5 agent steps	t_max	300 steps
Scout / firefighter view	radius 3 / 1	Water	10, 1 per cell
Stale after	15 steps	Scout target exclusion	5 cells
A* risk weight λ	5	Replan	every 10 steps or when blocked
Re-cluster and auction	every 5 steps	Two slots above	20 cells

Parameter	Value	Parameter	Value
Re-auction when away	> 15 steps	Fallback after	10 steps without AWARD
Score weights	0.7 / 0.2 / 0.1	Base station	(2, 2)

Results

Result	Value
Configurations where the auction beats independent significantly	7 of 8
Largest gain over independent	+25.3 points (10 firefighters)
Gain with 2 / 5 / 10 firefighters	+4.1 / +13.5 / +25.3 points
Auction vs greedy, 3 fires	+1.2 ± 0.9 points (significant)
River: auction better than independent	29 of 30 seeds
A* vs UCS	63% fewer nodes, same cost
Literal timing, all strategies	about 11.6% saved; perfect knowledge 15%
Utility targeting vs plan's downwind edge	+8.1 ± 4.2 points
Fallback vs no fallback after Coordinator failure	+16.8 ± 4.9 points
Firefighters caught per run (auction)	0.10
Tests / coverage	114 / 89%
Runs	720 main + 1,090 studies
Demo seed 0, windy: independent / greedy / auction	49.5% / 59.5% / 84.7% (contained at step 195)

Glossary

Term	Meaning here
Belief map	The team's record of every cell's last seen state and when it was seen
Stale	Seen more than 15 steps ago; scouts revisit it

Term	Meaning here
Frontier	Unknown or stale cells next to known, fresh ones
Zone	A connected group of burning cells auctioned as one task
Threat	Zone size $\times$ average chance its neighbouring trees ignite
Utility $U(f, z)$	Threat $\div$ (1 + path cost) $\times$ water fraction: a firefighter's bid
Slot	One firefighter position in a zone (1, or 2 above 20 cells)
Support round	Extra slots so no firefighter stays idle
Admissible heuristic	Never overestimates the remaining cost, so A* is optimal
Consistent heuristic	$h(n) \leq \text{step cost} + h(\text{next})$, so no node is expanded twice
Paired comparison	Strategies compared on the same seed, so the fire is identical
95% confidence interval	The range the true mean gain lies in with 95% confidence; not crossing 0 means significant
Fallback	Firefighter behaviour when no AWARD arrives: go to the nearest known fire
Oracle	Study-only mode where agents see the true fire state (an upper bound)
Contract net	The ANNOUNCE $\rightarrow$ BID $\rightarrow$ AWARD protocol for handing out tasks